

Introduction to Computational Linguistics

Session 5: Text Classification

Denise Löfflad

Universität Tübingen

November 29, 2023

- How computers learn
- Supervised and unsupervised machine learning
- Evaluating your model
- NLP segmentation

We are **not** going to talk about Naives Bayes, the Perceptron and other machine learning algorithms as this is for future semesters!

How Computers Learn

- machine learning: machine induces patterns from observations in data
- model:
 - summary of what the algorithm learned
 - estimates based on seen data, but can be applied to unseen data

How Computers Learn


Features

- Dimensions of interest in data
- Quantifications of raw data
- Measured for all records in the data set
- Individual features often combined into a *feature vector*
- Example: Speech Intelligibility scoring
 - Features:
 - rate of speech (syllables per second)
 - total time speaking (seconds)
 - amplitude (decibels)
 - pronunciation nativeness (GOP scores, Link)
 - number of grammatical errors (error count)
 - Feature Vector: $[1.7, 4, 55, -28, 2]$
- Learning algorithms accept feature vectors as input

How Computers Learn

data often split into different sets

- training set
- test set
- development/validation set or cross-validation



Source <https://www.researchgate.net/figure/>

The-learning-process-of-a-classifier-The-classifier-learns-from-the-training-data-The_fig1_333505093



Supervised Learning

- work with labelled data
- “supervised”: labels are determined by humans in the loop
 - e.g. annotations from corpora
- examples:
 - label how intelligible a recording is (good, fair, poor)

Supervised Learning

workflow:

- ① carefully formulate the RQ
- ② find or create appropriate data
- ③ define feature extracting mechanism and clean the data
- ④ decide on algorithm(s)
- ⑤ split the data
- ⑥ train the model
- ⑦ test and cross-validate
- ⑧ evaluate the model's performance
- ⑨ compare to baselines of StotA

compare with alternative algorithms/feature sets

Examples:

- linear regression
- logistic regression
- support vector machine
- decision tree
- classification tasks

Unsupervised Learning

- no pre-d categories
- purpose: automatically detecting structure in data

Unsupervised Learning

workflow:

- 1 define feature extracting mechanisms
- 2 decide on algorithms
- 3 apply on data set
- 4 inspect resulting inferred structure

Unsupervised Learning

examples:

- k-means clustering
- hierarchical clustering

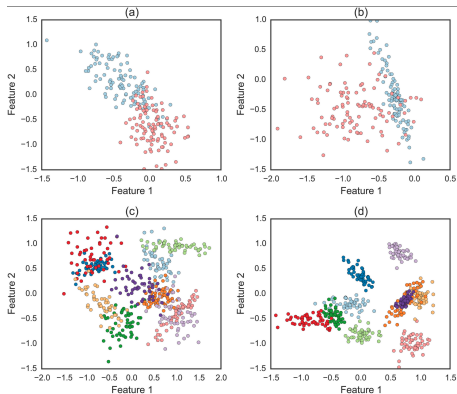


Figure: Taken from ,Rodriguez et al. [2019]

How to evaluate your (supervised) model?




Figure: Confusion matrices on classification data

Measuring Success (binary classification)

	Spam Classified	Ham Classified	
Spam	TP	FN	Recall
Ham	FP	TN	TNR
	Precision	True Omission Rate	

- Recall: $\frac{TP}{TP+FN}$
- Precision: $\frac{TP}{TP+FP}$
- Accuracy: $\frac{TN+TP}{TP+FP+FN+TN}$
- True negative rate (TNR): $\frac{TN}{FP+TN}$

Measuring Success

	Spam Classified	Ham Classified	
Spam	TP	FN	Recall
Ham	FP	TN	TNR
	Precision	True Omission Rate	

- Precision: $\frac{TP}{TP+FP}$

Measuring Success

	Spam Classified	Ham Classified	
Spam	TP	FN	Recall
Ham	FP	TN	TNR
	Precision	True Omission Rate	

- Recall: $\frac{TP}{TP+FN}$

Measuring Success - Error Analysis

	Spam Classified	Ham Classified	
Spam	TP	FN	Recall
Ham	FP	TN	TNR
	Precision	True Omission Rate	

- True negative rate (TNR): $\frac{TN}{FP+TN}$

Measuring Success


	Spam Classified	Ham Classified	
Spam	TP	FN	Recall
Ham	FP	TN	TNR
	Precision	True Omission Rate	

- Accuracy: $\frac{TN+TP}{TP+FP+FN+TN}$

Measures Success - Multiclass classification

For multiclass classification you have different options:

- calculate the values for every class
- average (micro or macro) over the classes



Which measures are important?

Which measures are important?

- Would you rather see spam mails or miss work mails?
- Would you rather see hate speech or get your normal Tweet deleted?

→ always consider the task you are working on! There is no one-fits-all solution.

NLP Tasks - Segmentation

- raw input: input stream of characters/text
- output: input segmented into smaller units
 - tokens/“words” → tokenization
 - sentence → sentence segmentation
 - chunks: groups of tokens

Tokenization

“A **token** is an *instance of a sequence of characters* in some particular document that *are grouped together as a useful semantic unit for processing*. A **type** is the class of all tokens containing the same character *sequence*. ”¹

¹<https://nlp.stanford.edu/IR-book/html/htmledition/tokenization-1.html>, last accessed 2021-01-18

Tokenization

- split into smaller units
- not necessarily always a “word”
- many decisions to make to deal with challenges

Tokenization – Challenges

- contracted/enclitic forms
- example: “I don’t know whether he’s going there.
- one/multiple words/tokens ?

Tokenization – Challenges

- hyphenated forms
- example: “An end user device-independent solution - would that be possible?”
- meaning preserved when splitting into two tokens?

Tokenization – Challenges

- forms with/adjacent to periods
- example:
“You find it on 21st St. opposite of Third Ave., in the U.S.A.”

Tokenization – Challenges

- slashes
- example:
“A helpful/fun e-learning website is <https://feedbook.schule>”

Tokenization – Challenges

- special characters
 - single/double quotation marks
 - comma/semicolon
 - question mark/exclamation marks
 - round brackets

- example:

“Yes!” is what I would ‘say’, but I don’t say it - I know better (now).

Tokenization – Challenges

- multi-word expressions
- includes compound nouns
- need to decide whether to recognize multiword tokens or not
- example: "hot dog", "in spite of"

Tokenization – Challenges

- named entities
- example: "New York"

Tokenization – Challenges

- abbreviations
- example: " Builders Inc. "

Tokenization – Challenges

- integrate morphological analysis to split certain contracted forms?
- usage of lexicon in tokenization?
 - Can help with things like “U.S.A.”, “hot dog”, and so on

Sentence Segmentation

- decide sentence boundaries
- in case of punctuation characters, decide whether they are
 - sentence-internal (e.g., Mr.)
 - sentence-final (e.g., The cat woke up.)
- can be tackled as classification task for punctuation
- many tools perform it in conjunction with tokenization

Further Challenges

- different meaning of characters in different language and writing systems
 - e.g. agglutinative languages vs. ideographic writing systems
- lack of punctuation or spaces in certain writing systems, changes in semantic interpretation depending on tokenization
- under-resourced languages: difficult to find corpora to train models

Getting Started with NLP - The OpenNLP Library

Basic properties of Apache **OpenNLP**:

- library with NLP tools for segmentation, tagging, named entity recognition, chunking, parsing, ...
- **machine learning** algorithms, some rule-based methods for comparison
- includes tools and pre-trained statistical language models
- accessible both as an **API (Application Programming Interface)** and as an executable program (**CLI** or **Command Line Interface**)

- official website: <https://opennlp.apache.org>
- OpenNLP Package by Chen Xiaobin:
<https://github.com/chxiaobin/opennlp/archive/master.zip>
- manuals at <https://opennlp.apache.org/docs/>

OpenNLP: Sentence Segmentation

Deciding whether a punctuation character marks the end of a sentence.

- `$ bin/opennlp SentenceDetector models/en-sent.bin < input/text1.txt`

Tokenization

Tokens: words, punctuation, numbers...

- tokenizer with statistical model
- `$ bin/opennlp TokenizerME models/en-token.bin < input/text1.txt`

Exercise

The OpenNLP library also includes a simple tokenizer which does not require a model for tokenization. Try to run it with the same input file and see if it gives you the result you want.

```
$ bin/opennlp SimpleTokenizer < input/text1.txt
```

References and Acknowledgments

These slides are largely based on Dickinson et al. [2012] and largely based on slides for Text Technology, 2023 by Stephen Bodnar and on Detmar Meurers' slides on Second Language Acquisition.

Markus Dickinson, Chris Brew, and Detmar Meurers. *Language and computers*. John Wiley & Sons, 2012.

Mayra Z Rodriguez, Cesar H Comin, Dalcimar Casanova, Odemir M Bruno, Diego R Amancio, Luciano da F Costa, and Francisco A Rodrigues. Clustering algorithms: A comparative approach. *PloS one*, 14(1):e0210236, 2019.